

Robust Low-Bitrate Neural Speech Coding for Resource-Constrained Communications

Andrew Thibodeau
Computer Science
Virginia Tech
Alexandria VA, USA
tandrew03@vt.edu

Wayne Cleary
Computer Science
Virginia Tech
Alexandria VA, USA
waynercleary@vt.edu

Matthew Hearne
Computer Science
Virginia Tech
Alexandria VA, USA

Abstract— Neural speech codecs offer the ability to deliver high-quality speech under a multitude of constraints that emulate real-world conditions. These speech codecs perform efficiently when constrained by compute power, memory, bitrate limits, and latency. Practical deployment on resource-constrained devices such as embedded systems and mobile devices requires simultaneous control of all constraints whilst accounting for acoustic conditions seen in the real world. This project investigates robust low-resource speech in the context of the IEEE LRAC 2.0 Challenge at ICASSP 2027. The project will begin with the provided baseline codec and expand to focus on noise and reverberation-aware training while adhering to the challenge constraints on bitrate, latency, and computation. The proposed solution aims to use a hybrid Linear Predictive Coding (LPC) - neural architecture. LPC analysis will separate the predictable speech from the excitation residual with a lightweight neural encoder and quantization stage will compress the residuals. The objective is to determine whether classical speech modeling can reduce the computational burden of learned components whilst preserving quality and intelligibility across multiple bitrate settings. The final system will be evaluated under the LRAC 2.0 constraints on latency, bitrate, and computational complexity using clean, noisy, reverberant, and multilingual speech. Experiments will compare the proposed LPC-neural architecture with the provided baseline and analyze quality-efficiency tradeoffs across 1, 6, 12, 24, and 32 kbps. Final evaluation will be based on a crowdsourced subjective battery with weights applied to different multilingual, goal, and test conditions.

Index Terms— hybrid architectures, low-bitrate speech coding, neural speech codecs, robust speech coding

I. INTRODUCTION

A. Motivation

There is an increasing need for speech interfaces on devices with limited processing power such as embedded systems and mobile devices, these devices require speech codecs that are efficient under different constraints [1]. Speech coding compresses the original signal into a compressed format for transmission and then is decoded to get the original speech

signal [2]. One such type of speech codec that has proved effective are neural speech codecs, these models are able to learn a mapping from raw audio signals to a compressed latent version and decode. The strength of these models is evident as they can achieve better perceptual quality at low bitrates but have higher latency when compared to other methods such as a traditional source-filter model [3]. Differing approaches to how a model process, transmits, and decodes audio proves to come with tradeoffs that makes deployment on resource-constrained devices in real-world environments an ongoing issue. The ability to develop a neural speech codec that is efficient in terms of compute power, memory, bitrate, and latency proves to be a difficult challenge.

B. Problem Statement

Many modern neural speech codecs learn use an end-to-end learned architecture which consists of a convolutional encoder, residual vector quantization, and a convolutional/GAN decoder pattern. These models prove to be accurate at reconstructing the original audio input, they require large architectures and higher latency which decrease model efficiency [3]. Hybrid approaches, ones which utilize other models' efficiency in combination with the reconstruction capabilities of neural speech codecs, have proved effective at increasing efficiency in real-world settings [4].

Linear Predictive Coding allows for a low computational load whilst being effective at low bitrates. LPC works by encoding short frames into parameters and decodes by reconstructing the source and filtering it through an estimated vocal tract model. This motivates a hybrid architecture in which a LPC model handles the predictable speech whilst a lightweight neural network can encode the information from the LPC's prediction residuals.

C. Significance

This project investigates if combining LPC with learned residuals from a lightweight neural network can combine the efficiency and quality aspects of the two individual approaches. We propose a hybrid PLC-Neural system in which LPC analysis separates speech into predictive parameters and an excitation residual, the neural network encoder and quantizer would then compress the residual

representation. The system will be evaluated against the LRAC 2.0 baseline across multiple bitrate settings, transparency, robustness to noise and reverb, and intelligibility through a crowdsourced subjective battery.

The central research question is whether explicitly removing predictable spectral structure through LPC allows a neural residual codec to achieve a better quality–complexity tradeoff than the LRAC 2.0 baseline while satisfying the challenge’s bitrate, latency, and compute constraints.

II. RELATED WORK

A. LPC

Linear predictive coding (LPC) is an all-pole source-filter model in which speech is modeled as the output. LPC is driven by an excitation signal and generates predictor coefficients which are estimated by minimizing mean-squares prediction error over short [5]. The way this model works is by decomposing the short-time spectrum into a flat excitation spectrum multiplied by a smooth spectral envelope which captures resonances [6]. LPC is primarily used for low-bit-rate speech, especially in the variant known as Codebook Excited Linear Speech (CELP).

B. End-to-End Neural Speech Codecs

There exists multiple end-to-end Neural Speech Codecs that perform all aspects of encoding, transmitting, and decoding. One such example would be that of SoundStream [7]. SoundStream labels itself as *a novel neural audio codec that can efficiently compress speech, music, and general audio at bitrates normally targeting by speech-tailored codecs*. The method in which the researchers accomplished this was by using an encoder/decoder network and a RVQ which were trained together end-to-end. The researchers demonstrated that the model could operate on bitrates from 3kbps to 18kbps with minimal loss, they also demonstrated the ability to run on a smartphone CPU indicating low latency. SoundStream shows that end-to-end neural speech codecs can operate at both a low bitrate and low latency, within the parameters of the LRAC 2.0 challenge, and perform in real world setting on devices such as mobile phones.

C. Hybrid Architectures

There exist hybrid approaches that combine the complexity of neural encoder/decoder networks as well as foundational speech models such as LPC. One of the most notable implementations is that of LPCNet [8]. The researchers utilized a WaveRNN variant that combines the linear prediction of LPC with RNN to improve the efficiency of speech synthesis. LPCNet uses LPC to represent the vocal-tract response and trains the neural component to predict the excitation residual. The motivation for this division of labor is that neural architecture can limit compute power as it doesn’t need to process the predictable component. The authors

conclude that their hybrid approach is acceptable for low bitrate speech coding.

The existing literature indicates that our proposed solution is a defensible approach as both the individual components, and the hybrid methodology, have been tested against similar metrics to those measured in the LRAC 2.0 challenge.

III. PROPOSED APPROACHES

We will first reproduce the LRAC 2.0 organizer baseline and establish reference benchmarks across the five (1, 6, 12, 24, 32) bitrates. This will provide us with the comparison point for evaluating whether the proposed LPC-neural hybrid architecture improves the quality-efficiency tradeoff.

The first aspect of implementing our proposed approach will consist of implementing a LPC front-end decomposition. This conceptually means that we will have our input audio waveform, and standard for how traditional LPC works will divide the waveform into short frames. On each frame, we will perform LPC analysis meaning that we will extract the LPC (known) coefficients of the speech, as well as compute the excitation residuals. The LPC coefficients derived in this step represent the predictable spectral structure of the speech whilst the residual information is what gets fed to the neural model. This breakdown allows for the neural model to deal with less data when compared to a compressed version of the entire waveform thus reducing the latency issue found with end-to-end neural model approaches.

The residual will then be passed to a lightweight neural encoder in which a compressed latent representation will be derived. From here, the compressed latent representation will then be quantized before having a lightweight neural decoder reconstruct the residual. After reconstruction, LPC synthesis will reconstruct the final speech.

After the architecture has been completed, we will train and evaluate with permitted noise augmentation, use room impulse responses for reverberation, and evaluate whether the LPC decomposition is effective with the degraded audio.

The final model will be compared with the organizer baseline as well as the initial model to see how the proposed architecture compares in both a clean setting as well as with degraded sound.

It is worth noting that although model selection is similar to that of other work such as LPCNet. LPCNet showed that linear prediction can reduce the complexity of neural speech synthesis via removing the predictable response as something the neural network models. What the researchers did was have the network predict the excitation signal from the LPC derived speech features which are combined with the linear prediction to generate the output. Our proposed system differs in both architecture and objective. We will use the LPC as a front-end decomposition, the LPC parameters represent the predictable spectral structure while the prediction residual is given to a

lightweight encoder, quantization stage, and neural decoder. This structure allows for us to evaluate bitrate, quantization, latency, and computational tradeoffs in accordance with the LRAC 2.0 constraints.

IV. SYSTEM DESIGN

The architecture for the proposed solution can be found in *Figure 1*.




Figure 1: Architecture

The proposed architecture will process a 24-kHz speech using frame-based LPC analysis, meaning the original speech is broken up into smaller chunks to be processed. The predictor coefficients will represent the short-term spectral envelope, while the prediction residual will be passed to a lightweight neural encoder. The encoder will produce a compressed latent representation that is then discretized via a quantization stage. A neural encoder will then reconstruct the residual signal before being combined with the LPC parameters with a LPC filter to reconstruct the waveform.

Due to the constraints of the LRAC 2.0 challenge, limited lookahead layers will be used to keep the latency of the system low (<30ms). The encoder and decoder will be benchmarked throughout development to measure MFLOPS, parameter count, and the overall complexity of the model. These steps would further simplify the complexity of the neural architecture past removing the predictable structure with the LPC front-end.

It is also worth noting that there are additional constraints laid out by the LRAC 2.0 Challenge, this means that our solution must have: total compute < 700 MFLOPS, total latency < 30ms, must support all bitrate caps and mixture of caps in one inference run, and track 2 requires < 1.5 million parameters.

The programming language that will be used is Python, this is chosen due to the available deep learning frameworks and associated libraries. We will utilize PyTorch as our deep-learning framework. For the audio processing, we will utilize torchaudio, NumPy, and SciPy. Visualization of benchmarking and intermediate testing will be done using Matplotlib. Development tools will consist of Git/GitHub, CUDA, and Virginia Tech resources including GLogin for computationally expensive workloads.

The datasets we will be using is the official LRAC 2.0 data provided on the challenge page. The dataset consists of permitted speech, noise, and RIR datasets that total 1200 hours of audio across 12 languages.

For evaluation of our proposed solution, we will evaluate over 4 different criteria. To measure quality, we will utilize the objective quality metrics provided. To measure intelligibility, we will compare with baseline metrics. We will also be looking at resource constraints by measuring bitrate, latency, MLFLOPS, parameter count, and memory cost. We will finally measure robustness by comparing clean speech, noisy speech, reverberant speech, and multilingual speech.

V. APPENDIX

The roles of each team member are as follows:

Name	Primary Role(s)	Responsibilities
Andrew	Team Lead, System Integration, and LRAC Compliance	Coordinate project milestones and team integration; maintain the main code repository; reproduce and integrate the LRAC baseline; verify bitrate, latency, and computational constraints; integrate LPC, neural codec, and evaluation components into the final system; coordinate proposal, presentations, and final report.
Matt	Data Pipeline, LPC Processing, and Robustness Augmentation	Prepare the LRAC 2.0 datasets using the official pipeline; implement frame-based LPC analysis and synthesis; generate LPC coefficients

		and excitation residuals; implement permitted noise and room-impulse-response augmentation; maintain clean, noisy, reverberant, and multilingual evaluation subsets.	October 14 – October 20 Integrate the initial lightweight neural residual encoder, quantization stage, and decoder. Perform preliminary baseline vs. LPC-neural comparisons. Prepare architecture and initial results. Checkpoint I due October 20.
			October 21 – October 31 Refine the LPC-neural architecture and required bitrate support. Implement permitted noise and reverberation augmentation. Begin clean vs. degraded-speech experiments.
Wayne	Neural Codec Development, Training, and Evaluation	Develop and train the lightweight neural residual encoder, quantization stage, and decoder; conduct baseline and proposed-model experiments; evaluate reconstruction quality and intelligibility; profile model size, MFLOPS, and latency; produce plots and experimental results for checkpoints and the final report.	November 1 – November 10 Conduct ablation experiments comparing the organizer baseline, LPC-neural codec, and robustness-enhanced LPC-neural codec. Evaluate multilingual performance, latency, MFLOPS, and quality-compute tradeoffs. Checkpoint II and 15-minute presentation due November 10.
			November 11 – November 21 Optimize the architecture based on checkpoint results. Investigate alternative quantization, model compression, or knowledge distillation if time permits. Complete major experiments and constraint-compliance testing.
			November 22 – November 30 Freeze the final model configuration. Complete rate-quality, robustness, multilingual, and efficiency experiments. Prepare final figures, tables, and LRAC submission materials.
			December 1 Process the LRAC blind-test data and submit challenge materials, if participating. Begin final project presentation period.
			December 1, 3, and 8 Final Project Presentations. Present final architecture, experimental results, ablation study, LRAC compliance, limitations, and conclusions on assigned date.
			December 2 – December 10 Incorporate presentation feedback, complete final analysis and discussion, clean repository/project webpage, and finalize the paper.
			December 11, 8:00 PM Final Report Due. Submit final report with methodology, implementation, experiments, results, analysis, limitations, and future work.

Although each team member has roles that they are responsible for, all members will contribute to each stage. The primary roles designation indicates which team member “owns” that step of the project.

The schedule is as follows:

Date / Period Planned Work and Deliverables	
September 22	Project Proposal Due. Finalize research question, LPC-neural architecture, related work, task assignments, schedule, and project webpage.
September 23 – October 3	Set up development environment and repository. Register for LRAC 2.0, obtain permitted datasets, run the official data-preparation pipeline, reproduce basic organizer baseline inference, and implement initial LPC analysis/synthesis.
October 4 – October 13	Reproduce and profile the organizer baseline. Measure bitrate, reconstruction quality, latency, computational complexity, and model size. Validate LPC decomposition and residual reconstruction on a development subset.

VI. REFERENCES

- [1] Cisco, “LRAC 2.0 Challenge,” *Low-Resource Audio Codec Challenge*, 2026. [Online]. Available: <https://crowdsourcing.cisco.com/lrac-challenge/2026/>. [Accessed: Sep. 22, 2026].
- [2] G. K. Sheferaw, W. Mwangi, M. W. Kimwele, and A. L. Mamuye, “Waveform based speech coding using nonlinear predictive techniques: A systematic review,” *International Journal of Speech Technology*, vol. 26, pp. 1031–1059, 2023, doi: 10.1007/s10772-023-10072-7.
- [3] X. Jiang, X. Peng, H. Xue, Y. Zhang, and Y. Lu, “Latent-domain predictive neural speech coding,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 31, pp. 2111–2123, 2023, doi: 10.1109/TASLP.2023.3277693.

- [4] M. Kim and J. Skoglund, "Neural speech and audio coding: Modern AI technology meets traditional codecs," *IEEE Signal Processing Magazine*, vol. 41, no. 6, pp. 85–93, Nov. 2024, doi: 10.1109/MSP.2024.3444318.
- [5] B. S. Atal and S. L. Hanauer, "Speech analysis and synthesis by linear prediction of the speech wave," *The Journal of the Acoustical Society of America*, vol. 50, no. 2B, pp. 637–655, 1971, doi: 10.1121/1.1912679.
- [6] J. O. Smith III, "Linear predictive coding of speech," in *Physical Audio Signal Processing*. Stanford, CA, USA: Stanford University, 2010. [Online]. Available: https://ccrma.stanford.edu/~jos/pasp/Linear_Predictive_Coding_Speech.html. [Accessed: Sep. 22, 2026].
- [7] N. Zeghidour, A. Luebs, A. Omran, J. Skoglund, and M. Tagliasacchi, "SoundStream: An end-to-end neural audio codec," *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 30, pp. 495–507, 2022, doi: 10.1109/TASLP.2021.3129994.
- [8] J.-M. Valin and J. Skoglund, "LPCNet: Improving neural speech synthesis through linear prediction," in *Proc. IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, Brighton, U.K., 2019, pp. 5891–5895, doi: 10.1109/ICASSP.2019.8682804.